

Orientation

what you have and how to run it

The BI theme digest agent, built for Momentive Software. This deck is for me, not for the review. It is the same forty five minute tour as docs/ORIENTATION.md, in slide form.

Three commands, twenty seconds, no network

```
make setup
make demo
open ../bi-theme-digest-store-demo/digests/2026-
W37.html
```

The demo clones the store's first commit into a scratch store, restores the recorded model responses, replays fifteen nightly ingests and three weekly builds, then runs the golden set. My last run finished in 19.5 seconds and printed 18 of 18 assertions passed and 19 commits.

In the HTML page, click an evidence line to expand it to the source moment, and click Why this score to expand the five weighted terms with the arithmetic shown.

```
# Theme digest, week 2026-W37
Run 2026-09-14T07:00Z, week 2026-W37: 25 sources read (14 calls, 11 cases), 4 private comments withheld, 5
  PII redactions, 27 claims verified, 0 rejected, 0 themes appended, 10 opened, cost $2.27 for the week
  ($1.22 this build), model tiers used: extraction, synthesis.
## Cold start: this week's claims resolve into ten themes, six of them carrying evidence from two separate
  customers.
The store started empty, so every claim opened a theme; the work was grouping, not triage. Six themes are
  backed by two customers each: truncated exports, wrong renewal amounts, pledge reminders that ignore donor
  preferences, renewal email landing in spam, duplicate entries from the accounting sync, and offline event
  check-in. Three of those six were only visible after reconciling different vocabularies, most notably one
  account's 'attendee kiosk' and another's 'event check-in', and Great Lakes' missing carried-over balance
  sitting on the same billing failure as Prairie Land's un-prorated mid-year upgrades. Four themes rest on a
  single voice: job posting expiry, tier migration tracking, an LMS SCORM import request from a prospect
  rather than a customer, and one piece of unprompted praise for the fundraising platform. Nothing is quiet
  or stale yet because nothing has history.
```

```
../bi-theme-digest-store/digests/2026-W37.md
```

The run line at the top of the markdown digest is the week in one sentence.

What lives in each top level folder

- `src/digest` is the package. Everything that runs is in there.
- `contracts` is the schema pack, and it is the law. Every boundary validates against a file in it.
- `mcp` holds the two mock servers, Gong and Salesforce, which enforce read only access at the boundary.
- `config` holds `models.yaml` and `models.cheap.yaml`, the only runtime files that name a model.
- `data/mock` holds the generated corpus, the four hand written traps, the accounts and the PII name list.
- `tests` holds 531 tests, one file per area, plus the fixtures.
- `evals` holds the golden set and its recorded results.
- `docs` holds decisions, governance, cost, production context, future builds, the verification guide and this orientation.
- `deck` holds the review deck, its outline and seventeen rendered panels.
- `tools` holds the corpus generator, `.claude` holds the same roles as subagents, `.github` holds the two crons.

Every file under src/digest, in five groups

Top level

pipeline.py runs a day, a week and the demo. router.py is the model seam. store.py is the only module that writes to the store. audit.py logs one event per action. contracts.py, errors.py, verify.py, enrich.py, score.py, eval.py, swap.py, gate.py and the CLI in __main__.py.

connectors

base.py is the protocol, gong.py and salesforce.py are the two implementations, mcp_client.py is a synchronous client over an asynchronous SDK.

agents

readers/reader.py plus two extraction prompts. editor/editor.py, tools.py and two prompts. analyst/analyst.py, retrieval.py and one prompt.

providers


agent_sdk_seat.py is the working default, anthropic_direct.py is implemented and unit tested, bedrock, vertex, foundry and openai_compatible are honest stubs with real model id translation. _families.py holds the per family thinking rules.

pii and render

pii/scrubber.py runs four regex passes before any model call. render holds markdown.py, html.py, citations.py and scoring.py.

Nine stages. Eight are code, one is judgment

1. Connect: connectors and mcp. No model.
2. Scrub: pii/scrubber.py. No model.
3. Extract: agents/readers/reader.py, extraction tier.
4. Validate and verify: verify.py. No model.
5. Enrich: enrich.py. No model.
6. Edit: agents/editor/editor.py, synthesis tier.
7. Score: score.py. No model.
8. Write and commit: store.py, render, audit.py.
9. Gate: gate.py. Nothing is filed without a person.



README.md

Off to the side, make ask runs the analyst on the narrative tier over the theme store only, with no access to the sources.

Which tier runs where, and where the prompt lives

extraction

claude-haiku-4-5

gong_reader.prompt.md and sfdc_reader.prompt.md, both at version 1.2.0. One scrubbed document in, one ReaderOutput out, no tools and no store.

synthesis

claude-opus-5

theme_request.prompt.md is call one and names the themes it needs. editor.prompt.md is call two and returns the proposal. Four read only tools and no write tool anywhere.

narrative

claude-opus-5

analyst.prompt.md. One question in, one grounded answer out, and it declines when the store does not support an answer.

No model id appears anywhere under src. Every call site asks for a tier, and config/models.yaml is the only runtime file that names a model.

Three layers, a watermark, and an evidence trail

ROUTER.md is layer 0, under five kilobytes, always read first. themes/_INDEX.md is layer 1 and is regenerated from the theme files on every write, so it cannot drift. themes/THEME-nnnn.md is layer 2, opened only when layer 1 says it is relevant. sources/_INDEX.md carries the watermark in its frontmatter, which is why a nightly run never re-reads the corpus.

```
themes/_INDEX.md, top 5 of 11 rows by score
```

```
THEME-0011 score 82 open integrations
Role changes are not propagated by the user sync, so people keep access they should have lost
aliases: identity provider sync; provisioning sync; role change; permission set
```

```
THEME-0002 score 81 open accounting
Renewal amounts ignore prior-period credits and mid-year upgrades, so members are over-billed
aliases: renewal statement; renewal invoice; partial year adjustment credit; mid-year upgrade
```

```
THEME-0001 score 75 open reporting
Report exports stop partway through above a row threshold, with no warning that data is missing
aliases: membership export; report export; year end donor report; data export
```

```
THEME-0006 score 75 open events
Event check-in cannot run without network and sync afterwards
aliases: event check-in; attendee kiosk
```

```
THEME-0003 score 70 open fundraising
Pledge reminders send on the standard schedule and channel regardless of donor preferences
aliases: pledge reminders; single reminder; mail only donors; paper only
```

```
../bi-theme-digest-store/themes/_INDEX.md
```

Eleven themes by week 39. Three quiet, none stale enough to drop, and the score column is computed in code.

The git log is the audit trail, one commit per run

The store is a separate repository with no code in it, so its history contains nothing but pipeline commits. Twenty six today: fifteen nightly ingests, three weekly builds, the scaffold, and the later ask and approve runs. Underneath that, runs/<run_id>/run.log.jsonl holds one validated event per action, reads included, and every scrub event carries counts and never values.

```
{"agent": "scrubber", "action": "scrub",
  "target": "500000000000000016",
  "outcome": "ok",
  "detail": {"EMAIL": 0, "PHONE": 0,
             "ADDRESS": 0, "NAME": 0}}
```

```
38204e1 run 2026-09-28T08:00Z: ask, answered
c237d18 run 2026-09-28T09:00Z: readme author name corrected
d4e5397 run 2026-09-28T09:00Z: approve dry run, nothing filed
65b7b0a run 2026-09-28T08:00Z: ask, declined
212fc38 run 2026-09-28T08:00Z: ask, answered
bf59c93 run 2026-09-14T07:00Z: swap 2026-W37 onto models.cheap.yaml, jaccard 1.000
e93f5ba run 2026-09-14T07:00Z: stability for 2026-W37, jaccard 1.000
05ec3a2 run 2026-09-28T07:00Z: build 2026-W39 digest, 8 themes appended, 0 opened
99b01b1 run 2026-09-21T07:00Z: build 2026-W38 digest, 7 themes appended, 1 opened
c51b740 run 2026-09-14T07:00Z: build 2026-W37 digest, 0 themes appended, 10 opened
ebddf72 run 2026-09-25T06:00Z: ingest 2026-09-25, 5 sources, 5 claims
8c33228 run 2026-09-24T06:00Z: ingest 2026-09-24, 1 sources, 1 claims
cf44727 run 2026-09-23T06:00Z: ingest 2026-09-23, 3 sources, 3 claims
5582b25 run 2026-09-22T06:00Z: ingest 2026-09-22, 2 sources, 2 claims
7f75cee run 2026-09-21T06:00Z: ingest 2026-09-21, 1 sources, 1 claims
57df964 run 2026-09-18T06:00Z: ingest 2026-09-18, 2 sources, 2 claims
578d23f run 2026-09-17T06:00Z: ingest 2026-09-17, 4 sources, 4 claims
e772967 run 2026-09-16T06:00Z: ingest 2026-09-16, 4 sources, 4 claims
0cc12d1 run 2026-09-15T06:00Z: ingest 2026-09-15, 3 sources, 3 claims
44428f6 run 2026-09-14T06:00Z: ingest 2026-09-14, 3 sources, 3 claims
7919d63 run 2026-09-11T06:00Z: ingest 2026-09-11, 4 sources, 4 claims
0f1a3f5 run 2026-09-10T06:00Z: ingest 2026-09-10, 5 sources, 5 claims
dcf1a80 run 2026-09-09T06:00Z: ingest 2026-09-09, 7 sources, 7 claims
ccffc89 run 2026-09-08T06:00Z: ingest 2026-09-08, 9 sources, 11 claims
896dde5 run 2026-09-07T06:00Z: ingest 2026-09-07, 0 sources, 0 claims
8d84584 scaffold: empty context store, ROUTER.md and the two index files
```

```
$ git -C ../bi-theme-digest-store log --oneline
```


What each edit invalidates

Change a prompt

Bump prompt_version in the frontmatter. The replay key is a hash of tier, system, user and schema name, so editing a prompt invalidates every recording for that agent and make demo stops with exit 3, a replay miss. It never falls back to a live call. Regenerating means a record run, about 54 minutes and real money.

Change a tier

Edit config/models.yaml. The replay key uses the tier name and not the model id, so replay will hand back the old model's answer. To compare models properly, run make swap-models, which records the alternate side into its own directory.

Add a source

One file in src/digest/connectors implementing list_since and fetch, plus a config entry. The pipeline does not grow a branch.

Run one test file

The venv interpreter with pytest and the file name. tests/test_score.py passes eleven tests in 0.15 seconds. make test runs all 531.

make swap-models, on the cheap tier map

```
cost:    reference $1.2187, alternate $0.2831 (23.2 percent of reference)
eval:    reference 16 of 16 assertions, alternate 16 of 16
overlap: claim co-assignment jaccard 1.000 (with theme ids 0.174), 10 themes against 10, top three CHANGED
        (['THEME-0002', 'THEME-0001', 'THEME-0006'] vs ['THEME-0007', 'THEME-0001', 'THEME-0003'])
```

```
$ make swap-models # python -m digest swap --models config/models.cheap.yaml --week 2026-W37
```

Twenty three percent of the reference cost, and it grouped every claim the same way. The theme ids on the top three differ because ids are handed out in the editor's placeholder order, which on a cold start week measures the labelling and not the ranking.

Six waves, thirty briefs, one envelope each

6

waves under one orchestrator,
from the contract pack to the
deck

30

briefs, each owning an exclusive
set of file paths

30

JSON envelopes returned, every
one with an assumptions array

531

tests, written by the worker that
owned the area

Wave 0 was the contract pack on its own, because fourteen later workers had to build in parallel without talking to each other. Waves 1 and 2 were seven workers each, wave 3 was the pipeline and the first live run, wave 4 the documents, wave 5 the QA auditors, wave 6 the deck.

Every worker read only the contract slices its brief named, wrote only inside its own paths, and returned a structured envelope rather than prose. Anything it had to infer went into a named assumptions array instead of being buried in code. Context rot is the failure mode of a large agent build, and the fix is narrow briefs and exclusive ownership rather than a bigger context window.

Run one to run three, and what the eval caught

Run 1

83 claims, 20 themes against a plan of nine, eval 17 of 18, a week 37 rerun agreeing with itself at 0.79. Measured and not shipped. Three causes, each fixed at its owner: filler in the corpus carried extractable asks, the readers split one point into two claims, the editor split facets of one problem into themes.

Run 2

55 claims, 12 themes, rerun agreement 1.000. The golden set still caught one reader recall miss, so the reader prompts went to 1.2.0 with a recall first rule and a stricter praise test.

Run 3

55 claims from 53 sources, 11 themes, 18 of 18, zero rejected, 5.47 dollars by the rate table. This is the run that is committed and the one the replay reproduces byte for byte.

```
evals/golden_set.yaml via evals/results.md: 5 must-appear, 2 must-not-appear
```

```
T1a [pass] The renewal invoice claim from the Gong call 7782934451002 appears in the store, cited,
          attached to a theme, and quoted in the digest.
          evidence: theme THEME-0002, claim 03a7ca4ae6d8
T1b [pass] The renewal invoice claim from the Salesforce case 5008W00002aQpLrQAK appears in the store,
          cited, attached to a theme, and quoted in the digest.
          evidence: theme THEME-0002, claim 306b8bbab997
T2a [pass] The event check-in claim from the Gong call 7782934451118 appears in the store, cited,
          attached to a theme, and quoted in the digest.
          evidence: theme THEME-0006, claim dcb169d12187
T2b [pass] The attendee kiosk claim from the Gong call 7782934451207 appears in the store, cited,
          attached to a theme, and quoted in the digest.
          evidence: theme THEME-0006, claim ae96fcdf3b0a
T3 [pass] The SCORM upload claim from the Gong call 7782934451311 appears in the store, cited, attached
          to a theme, and quoted in the digest, even though the account is a prospect.
          evidence: theme THEME-0009, claim 61d958ff323f
T4a [pass] The private Salesforce case comment about churn risk never reaches any output.
          evidence: absent from every file in the store
T4c [pass] The Momentive Software account executive line never becomes a client claim in any output.
          evidence: absent from every output; present only in the files that hold what the reader was shown
```

```
Result: 18 of 18 assertions passed, exit code 0.
```

```
evals/results.md
```

The golden set is the thing that caught the reader miss twice. It is mechanical: joins, substring tests and comparisons, and it asks no model anything.

The QA wave built nothing, and each caught something

Acceptance

Found make demo failing on a fresh clone, because the weekly build was not idempotent and rebuilt over its own output. That produced the already built guard and the scratch store demo. The second audit reads 21 pass, 0 fail, 3 untestable here.

Citation

Resolved all 55 claims and all 125 digest evidence lines by hand with its own throwaway resolver, importing nothing from the package. It caught the doubled PII redaction count, which was the one real bug.

Compliance

Scanned both repositories for dashes, non ASCII, model ids outside config, URL shaped text and voice drift. It found a local home directory path leaking into evals/results.md.

Deck

Traced every number on every slide back to a manifest field and failed the deck on a stale whole store total, with a numbered fix list. Those fixes are in.

Whoever wrote a piece never checked it. That two agent gate is what found the redaction bug and the fresh clone bug, and neither one was caught by the pipeline's own tests.

The things I would rather say first than be asked

- One sentence in the week 37 digest calls a claim the one medium importance claim in the set, when there are two. I left it rather than hand edit a recorded run's output.
- Week 37 opened ten themes where the plan expected nine. The extra one is real and cited, and tuning it away would have been tuning to my plan rather than to the evidence.
- The stability number is honest and awkward. Two live opinions grouped every claim identically, but the older flag compared theme ids and read false. Both numbers are printed.
- Three acceptance lines are untestable locally: live mode, a real issue being filed, and the deck audit.
- The citation auditor could not reproduce the dollar figures by hand, because nothing in the store carries a price.
- evals/golden_set.yaml carries the raw planted PII values in a file that is going public. Functionally correct, worth a second opinion.
- The manifest counts propose events, so the weekly manifests read 13, 11 and 9 while the proposal directories hold 6, 5 and 4. Both numbers are published.
- make approve already carries the yes flag and defaults to a real repository. Never run it to show what the gate does. Use the dry run, or the bare command that exits 4.

Eight beats, in this order

1

make demo

Fifteen ingests and three builds, twenty seconds, zero network, byte identical to what is committed.

2

open the week 37 markdown digest

Read the run line, then theme 1 and its three evidence lines.

3

grep the claim id in the evidence file

Point at the source reference: call, speaker, window, and the side derived by code.

4

grep the sentence in the raw call file

Say the rule once. The quote is an exact substring of that window or the claim is rejected.

5

grep the withhold action in a run log

The withholding happened in the query. The log carries the count and never the value.

6

grep the editor tool names

Four reads. A grep for a write function matches nothing.

7

approve, dry run, then without the yes flag

It prints the issue and files nothing, then exits 4, refused.

8

swap models, and one unanswerable question

Twenty three percent of the cost with the same groupings, then a visible decline.

How to record it, and what to verify first

Recording the walkthrough

Terminal at a large font, nothing else on screen, browser ready on the HTML digest in a second tab. Run make demo once before recording so the caches are warm and the scratch store already exists. Do not narrate the setup. Start with the first command typed and unexecuted, keep to the eight beats, and take it in one pass. Afterwards, put the recording location into the line at the top of walkthrough.md.

Before Momentive Software tests it

- make test reads 531 passed.
- make demo from a clean tree reads 18 of 18 and 19 commits.
- git status is empty in both repositories.
- The store is at its committed state, or make clean-store puts it back.
- The dash grep over both repositories prints nothing.
- The home directory path in evals/results.md is gone before these go public.
- docs/VERIFY.md is the page I send. It tells a stranger what to type and what should come back.